


```
In [1]: import tensorflow
        from tensorflow import keras
        import matplotlib.pyplot as plt
```

```
/Users/lufy/Developer/projects/digit-recognizer/venv/lib/python3.9/site-packages/urllib3/__init__.py:35: NotOpenSSLWarning: urllib3 v2 only supports OpenSSL 1.1.1+, currently the 'ssl' module is compiled with 'LibreSSL 2.8.3'. See: https://github.com/urllib3/urllib3/issues/3020
warnings.warn(
```

```
In [2]: (x_train, y_train), (x_test, y_test) = keras.datasets.mnist.load_data()
```

```
In [ ]: # Training image
        plt.imshow(x_train[0], cmap="gray")
```

```
Out[ ]: <matplotlib.image.AxesImage at 0x1586a7490>
```



```
In [4]: # Pixel value harulai [0, 1] ko range ma lyaune (normalization)
        x_train = x_train / 255
        x_test = x_test / 255
```

```
In [5]: # Create keras sequential model
        model = keras.Sequential()

        # First layer of the neural network, this is where the image data goes.
        # The image is 28*28 pixels, so we need to flatten the image into 784 val
        # 784 nodes hune vayo yo layer ma
        model.add(keras.layers.Flatten(input_shape=(28,28)))

        # Second layer of the neural network, this layer is hidden and contains 1
        model.add(keras.layers.Dense(128, activation="relu"))

        # Last layer containing 10 nodes. 10 ota number 0, 1, 2, .. 9 vako le
```

```
# 10 ota final nodes
# classification ko lagi "softmax" use garne
model.add(keras.layers.Dense(10, activation="softmax"))
```

/Users/lufy/Developer/projects/digit-recognizer/venv/lib/python3.9/site-packages/keras/src/layers/reshaping/flatten.py:37: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.
 super().__init__(**kwargs)

In [6]: model.summary()

Model: "sequential"

Layer (type)	Output Shape	
flatten (Flatten)	(None, 784)	
dense (Dense)	(None, 128)	
dense_1 (Dense)	(None, 10)	


Total params: 101,770 (397.54 KB)


Trainable params: 101,770 (397.54 KB)


Non-trainable params: 0 (0.00 B)


In [7]: # Model lai compile hanne, ani train garna milxa
 model.compile(loss=keras.losses.SparseCategoricalCrossentropy, optimizer=


In [8]: # Model lai train garne, epochs = num of times to train
 # validation split = 0.2 vaneko purai data ma 20% chai
 # validation garna ko lagi chuttyayeko
 history = model.fit(x_train, y_train, epochs=25, validation_split=0.2)


Epoch 1/25
1500/1500  2s 1ms/step - accuracy: 0.8647 - loss: 0.48
12 - val_accuracy: 0.9542 - val_loss: 0.1628


Epoch 2/25
1500/1500  2s 1ms/step - accuracy: 0.9610 - loss: 0.13
19 - val_accuracy: 0.9650 - val_loss: 0.1163


Epoch 3/25
1500/1500  2s 1ms/step - accuracy: 0.9737 - loss: 0.08
96 - val_accuracy: 0.9704 - val_loss: 0.0989


Epoch 4/25
1500/1500  2s 1ms/step - accuracy: 0.9817 - loss: 0.06
36 - val_accuracy: 0.9722 - val_loss: 0.0974


Epoch 5/25
1500/1500  2s 1ms/step - accuracy: 0.9850 - loss: 0.04
95 - val_accuracy: 0.9744 - val_loss: 0.0904


Epoch 6/25
1500/1500  2s 1ms/step - accuracy: 0.9894 - loss: 0.03
88 - val_accuracy: 0.9750 - val_loss: 0.0894


Epoch 7/25
1500/1500  2s 1ms/step - accuracy: 0.9919 - loss: 0.02
86 - val_accuracy: 0.9751 - val_loss: 0.0863


Epoch 8/25
1500/1500  2s 1ms/step - accuracy: 0.9929 - loss: 0.02
39 - val_accuracy: 0.9745 - val_loss: 0.0932


Epoch 9/25
1500/1500  2s 1ms/step - accuracy: 0.9954 - loss: 0.01
67 - val_accuracy: 0.9700 - val_loss: 0.1135


Epoch 10/25
1500/1500  2s 1ms/step - accuracy: 0.9955 - loss: 0.01
49 - val_accuracy: 0.9749 - val_loss: 0.1052


Epoch 11/25
1500/1500  2s 1ms/step - accuracy: 0.9964 - loss: 0.01
21 - val_accuracy: 0.9757 - val_loss: 0.0940


Epoch 12/25
1500/1500  2s 1ms/step - accuracy: 0.9977 - loss: 0.00
97 - val_accuracy: 0.9765 - val_loss: 0.1024


Epoch 13/25
1500/1500  2s 1ms/step - accuracy: 0.9973 - loss: 0.00
94 - val_accuracy: 0.9759 - val_loss: 0.1074


Epoch 14/25
1500/1500  2s 1ms/step - accuracy: 0.9982 - loss: 0.00
73 - val_accuracy: 0.9762 - val_loss: 0.1094


Epoch 15/25
1500/1500  2s 1ms/step - accuracy: 0.9983 - loss: 0.00
70 - val_accuracy: 0.9729 - val_loss: 0.1184






Epoch 16/25
1500/1500  2s 1ms/step - accuracy: 0.9980 - loss: 0.00
71 - val_accuracy: 0.9758 - val_loss: 0.1172

Epoch 17/25
1500/1500  2s 1ms/step - accuracy: 0.9983 - loss: 0.00
60 - val_accuracy: 0.9747 - val_loss: 0.1197


Epoch 18/25
1500/1500  2s 1ms/step - accuracy: 0.9980 - loss: 0.00
67 - val_accuracy: 0.9762 - val_loss: 0.1119

Epoch 19/25
1500/1500  2s 1ms/step - accuracy: 0.9986 - loss: 0.00
50 - val_accuracy: 0.9777 - val_loss: 0.1146

Epoch 20/25
1500/1500  2s 1ms/step - accuracy: 0.9995 - loss: 0.00
28 - val_accuracy: 0.9725 - val_loss: 0.1460

Epoch 21/25
1500/1500  2s 1ms/step - accuracy: 0.9991 - loss: 0.00
 41 - val_accuracy: 0.9774 - val_loss: 0.1230
 Epoch 22/25
1500/1500  2s 1ms/step - accuracy: 0.9994 - loss: 0.00
 27 - val_accuracy: 0.9746 - val_loss: 0.1394
 Epoch 23/25
1500/1500  2s 1ms/step - accuracy: 0.9978 - loss: 0.00
 61 - val_accuracy: 0.9731 - val_loss: 0.1476
 Epoch 24/25
1500/1500  2s 1ms/step - accuracy: 0.9985 - loss: 0.00
 47 - val_accuracy: 0.9753 - val_loss: 0.1432
 Epoch 25/25
1500/1500  2s 1ms/step - accuracy: 0.9996 - loss: 0.00
 19 - val_accuracy: 0.9725 - val_loss: 0.1574

```
In [9]: # Predict the test inputs
y_probabilities = model.predict(x_test)
y_predictions = y_probabilities.argmax(axis=1)
print("test y values:", y_test)
print("test y predictions:", y_predictions)
```

313/313  0s 408us/step
 test y values: [7 2 1 ... 4 5 6]
 test y predictions: [7 2 1 ... 4 5 6]


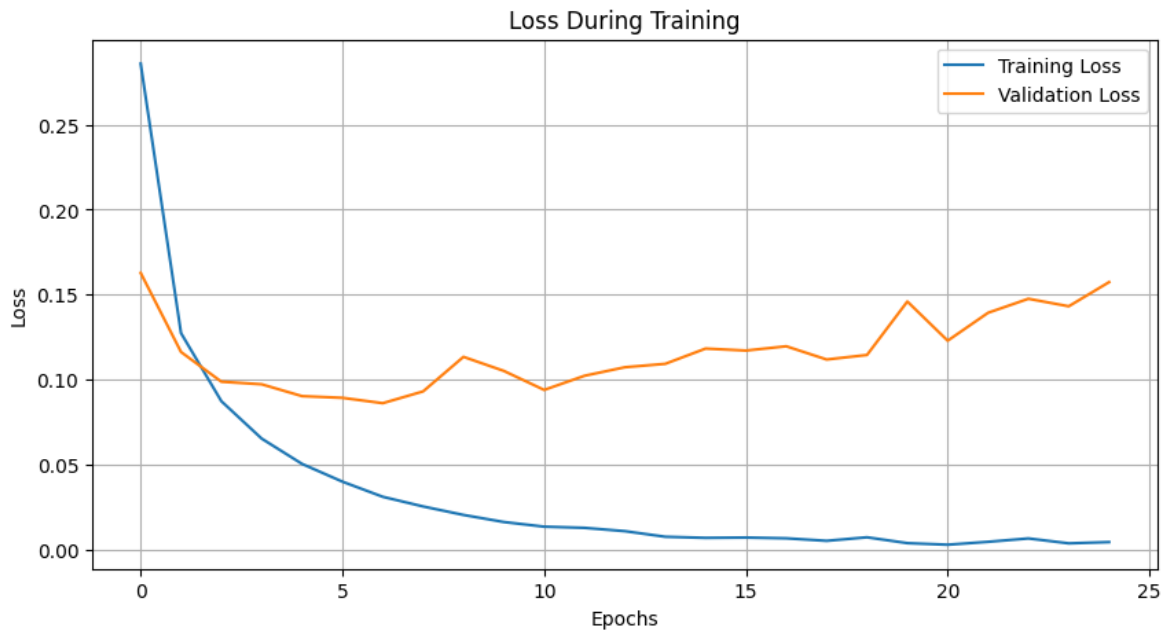
```
In [10]: # Aha accuracy measure karne model ko
# accuracy improve karne layer badhane milo, epoch value abhi dherai hai
# overfitting huna sakta hai
from sklearn.metrics import accuracy_score
accuracy_score(y_test, y_predictions)
```

Out[10]: 0.9748

```
In [18]: # Analyze the training process

# Plot training and validation loss
plt.figure(figsize=(10, 5)) # figure wide banana
plt.plot(history.history["loss"], label="Training Loss")
plt.plot(history.history["val_loss"], label="Validation Loss")
plt.title("Loss During Training")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()
plt.grid()
plt.show()

# Plot training and validation accuracy
plt.figure(figsize=(10, 5))
plt.plot(history.history["accuracy"], label="Training Accuracy")
plt.plot(history.history["val_accuracy"], label="Validation Accuracy")
plt.title("Accuracy During Training")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend()
plt.grid()
plt.show()
```




```
In [ ]: # Test our model's output with 10 random test inputs
from random import randrange

for i in range(10):
    input_image = x_test[randrange(len(x_test))]
    plt.imshow(input_image, cmap="gray")
    probabilities = model.predict(input_image.reshape(1, 28, 28))
    prediction = probabilities.argmax(axis=1)[0] # one with highest proba
    plt.title(f"Prediction: {prediction}")
    plt.axis('off')
    plt.show()
```


1/1 ————— 0s 19ms/step

Prediction: 1




1/1 ————— 0s 18ms/step

Prediction: 3




1/1 ————— 0s 18ms/step

Prediction: 4




1/1 ————— 0s 18ms/step

Prediction: 4




1/1 ————— 0s 18ms/step

Prediction: 6




1/1 ————— 0s 19ms/step

Prediction: 8




1/1 ————— 0s 18ms/step

Prediction: 6




1/1 ————— 0s 49ms/step

Prediction: 2



1/1 ————— 0s 16ms/step

Prediction: 9



1/1 — 0s 19ms/step

Prediction: 0

